Vidit Dhull, Data Engineer

vidit.dh@gmail.com | Portfolio: viditdhull.github.io | +91 9355512200 | LinkedIn: linkedin/viditdhull/ | Bengaluru, India

PROFESSIONAL SUMMARY

- **Data Engineer** with **5+ years** of experience building **scalable data platforms** and **real-time streaming pipelines** using Spark, Kafka, and both **open-source** and **Azure-based data ecosystems**.
- Experienced in designing Medallion Delta Lakehouse architectures, optimizing large-scale ETL pipelines, and delivering **high-throughput event-driven systems** processing TB-scale telemetry data.
- **Strong ownership mindset** with hands-on **leadership in data platform architecture**, performance optimization, and reusable engineering standards. Developed **AI-assisted SDLC automation** using Azure OpenAI, agentic workflows, and developer productivity tooling.

WORK AUTHORIZATION

- Non-EU. Requires employer sponsorship (Highly Skilled Migrant / EU Blue Card). Ready to relocate to Netherlands.
- **Languages: English** - Fluent/Professional; **Hindi** - Native; **Dutch** - A1 (learning)

EXPERIENCE

Data Engineer, Alstom **March 2023 - Present**

- Took **Dialogs**, a company-wide log analytics platform for train telemetry subsystems from POC to production on a **Medallion lakehouse**; unified heterogeneous logs and delivered **Power BI dashboards** for KPI monitoring, log exploration, and pattern analysis; **reduced recurring analysis cycles from 2 months to under 1 week**.
- **Owned end-to-end** delivery of an **auto-trigger, machine-to-machine ingestion** framework in **Microsoft Fabric** (Blob → Event stream/hub → Lakehouse/Warehouse/Eventhouse), supporting **30+ GB (5M+ events)** daily data ingestion.
- **Led architecture + implementation** of **signature/pattern search** on telemetry using **PySpark**, enabling fast matching of event sequences to known signatures; reduced investigation time by **50%** for recurring issues.
- **Designed** and **delivered** partial-run handling (late/partial/complete arrivals) with **idempotent merges** + governance rules, cutting data inconsistencies and manual rework by **60%** and improving KPI reliability.
- **Improved performance and cost-efficiency** by optimizing **Spark workloads** (partitioning, broadcast joins, shuffle reduction, handling data skew); cut backfill time by **50%** and improved query latency by **90%**.
- Implemented **agentic automation** for Microsoft Fabric pipelines using **Azure AI Foundry**, **MCP-style tool integration**, to support data quality checks, test generation, validation, spec-driven development.
- **Technical owner** for the CCP2 subsystem, **leading 3 data engineers**, translating client **requirements into scalable solutions**, assigning implementation tasks, reviewing code, and **driving end-to-end delivery** of features/data**.**

Software Engineer, Nextsavy Technologies **September 2020 - August 2022**

- Built **Azure Databricks pipelines** for mobility data using **ADF, PySpark, Delta Lake**, incremental loads, schema enforcement, and bad-record quarantine, **reducing reprocessing overhead by 60%**.
- Developed **Snowflake dimensional marts** with fact/dimension modeling, **SCD handling**, clustering, RBAC, and quality checks, **improving BI query performance by 50%** for revenue, demand, and utilization reporting.

SKILLS

- **Languages:** Python, SQL, PySpark
- **Data Engineering:** Apache Spark, Apache Airflow, Kafka, Batch and Real-time processing, Delta/Parquet
- **Cloud/Data Platforms:** Microsoft Fabric, ADLS Gen2, ADF, Azure Synapse, Azure Functions, Databricks, Snowflake, OneLake and Delta Lake concepts, Azure Storage/Blob, Events stream, Power BI, KQL/Eventhouse
- **Data Modeling:** Medallion/Lakehouse Architecture, Star Schema, ELT/ETL, Indexing, Partition, Dimensional Modeling
- **Data Quality & Governance:** Validation Automation, Logging Frameworks, KPI Modeling, RBAC, Access Control
- **Collaboration & Version Control:** GitLab, Agile/Scrum, CI/CD
- **Additional Skills:** Requirement Analysis, FinOps Awareness, Code Reviews, Stakeholder Collaboration

CERTIFICATIONS

- **DP-700:** Fabric Data Engineer Associate (**Microsoft**), September 2025
- **DP-203:** Azure Data Engineer Associate (**Microsoft**), June 2024
- **DP-900:** Azure Data Fundamentals (**Microsoft**), April 2024
- **IBM:** Data Engineering Specialization, October 2023

EDUCATION

- **PG-Diploma, Big Data Analytics (Ranked 1st)** - CDAC Bengaluru | (September 2022 - March 2023)
- **B. Tech in Computer Science** - UIET-MDU, Haryana | (July 2018 - July 2022)